

Evaluating Classifiers

Machine Learning

PHYS 453 – Spring 2024

Dr. Daugherty

<https://www.youtube.com/watch?v=--s zrOHtR6U>



VeggieTales: Monkey Silly Song



VeggieTales Official ✓
567K subscribers

Subscribe

👍 14K



➦ Share

⌵ Save



Evaluating Classifiers

- How can I measure how well a classifier works?
- Where do I look for ways to improve performance?

Sources:

- https://scikit-learn.org/stable/modules/model_evaluation.html#classification-metrics
- Binary Classification Metrics paper, on canvas or: <https://arxiv.org/pdf/1410.5330>
- https://scikit-learn.org/stable/auto_examples/model_selection/plot_precision_recall.html

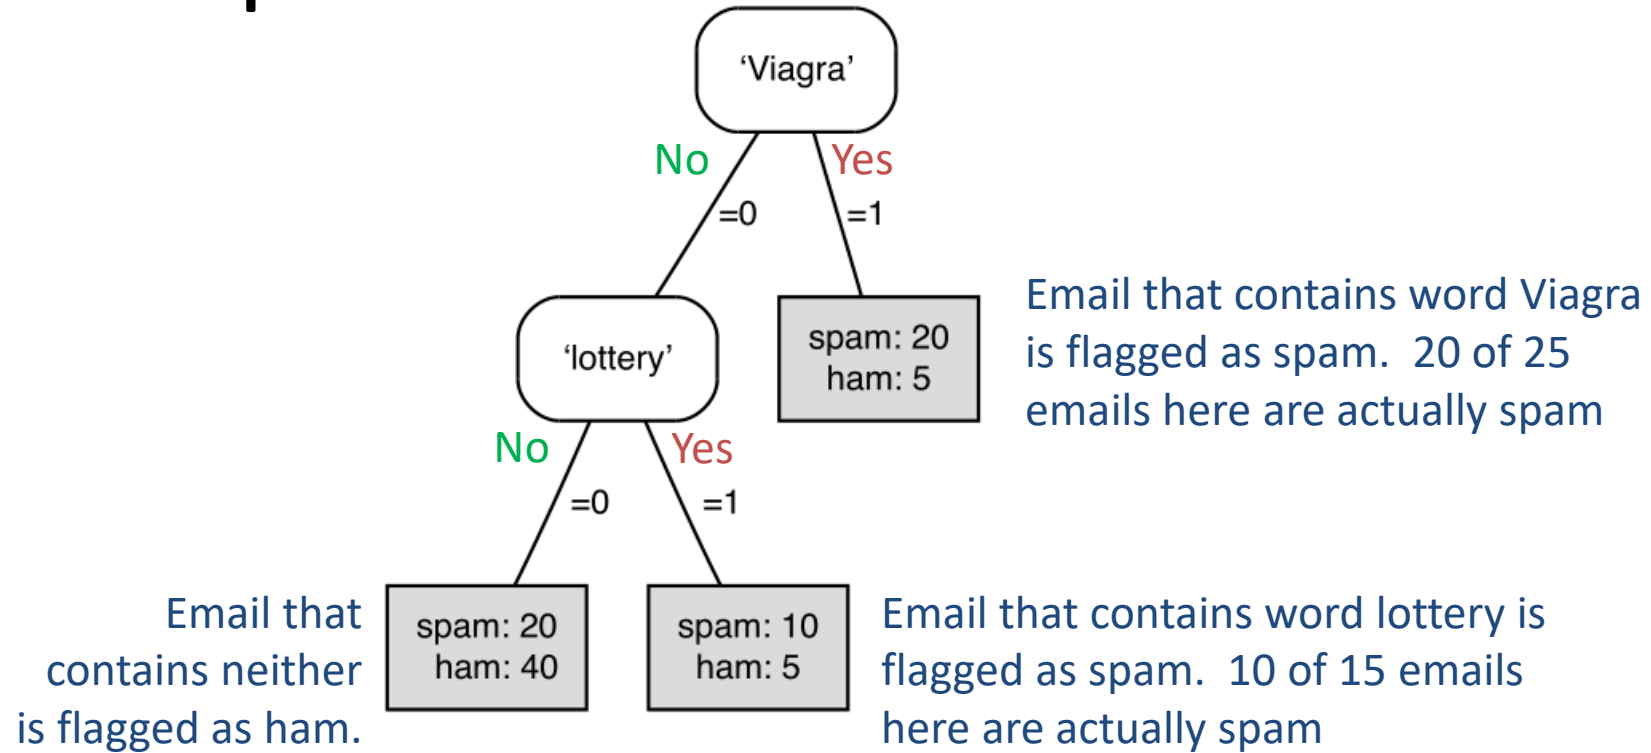
Evaluating Classifiers

CONFUSION MATRIX

Confusion Matrix

		Predicted class	
		P	N
Actual Class	P	True Positives (TP)	False Negatives (FN)
	N	False Positives (FP)	True Negatives (TN)

Spam detection



		Predicted class	
		<i>P</i>	<i>N</i>
Actual Class	<i>P</i>	True Positives (TP)	False Negatives (FN)
	<i>N</i>	False Positives (FP)	True Negatives (TN)

	SPAM		HAM	
	<i>Predicted</i> \oplus		<i>Predicted</i> \ominus	
<i>Actual</i> \oplus	30	20	50	
<i>Actual</i> \ominus	10	40	50	
	40	60	100	


Chapter 2.1

Difficulty: 2

There are 20 dogs(+) and 10 cats(-). A binary classifier correctly predicts 5 dogs and incorrectly predicts 5 cats. Fill in the following contingency for this binary classifier matrix.

	Predicted +	Predicted -	
Actual +			
Actual -			

Shows the biggest source of error is calling everything an 8 (especially 3's)



Evaluating Classifiers

CLASSIFIER METRICS

<https://xkcd.com/2236/>

IS IT CHRISTMAS?

|<

< PREV

RANDOM

NEXT >

>|

NO*

*99.73% ACCURATE

XKCD.COM PRESENTS A NEW "IS IT CHRISTMAS"
SERVICE TO COMPETE WITH ISITCHRISTMAS.COM

|<

< PREV

RANDOM

NEXT >

>|

Accuracy Metrics

		Predicted class	
		<i>P</i>	<i>N</i>
Actual Class	<i>P</i>	True Positives (TP)	False Negatives (FN)
	<i>N</i>	False Positives (FP)	True Negatives (TN)

$$ERR = \frac{FP + FN}{FP + FN + TP + TN} = 1 - ACC$$

Error %

$$ACC = \frac{TP + TN}{FP + FN + TP + TN} = 1 - ERR$$

Accuracy %

$$FPR = \frac{FP}{N} = \frac{FP}{FP + TN}$$

False Positive Rate = (# of FP) / (# actually N)
“what percentage of the real N did I miss?”

$$TPR = \frac{TP}{P} = \frac{TP}{FN + TP}$$

True Positive Rate = (# of TP) / (# actually P)
“what percentage of the real P did I get?”

$$PRE = \frac{TP}{TP + FP}$$

PRECISION = the ability of the classifier not to label as positive a sample that is negative.
Fraction of pos guesses that are right.

$$REC = TPR = \frac{TP}{P} = \frac{TP}{FN + TP}$$

RECALL (aka TPR) = the ability of the classifier to find all the positive samples
Fraction of all actual pos we guessed as pos.

$$F_1 = 2 \cdot \frac{PRE \cdot REC}{PRE + REC}$$

F1 Score = combines both into a single number. 1 is perfect.

Metrics

- A system with high recall but low precision returns many results, but most of its predicted labels are incorrect when compared to the training labels.
- A system with high precision but low recall is just the opposite, returning very few results, but most of its predicted labels are correct when compared to the training labels.
- Can plot curve showing precision/recall trade-off

Challenge: gotta find them all!

	Predicted \oplus	Predicted \ominus	
Actual \oplus	30	20	50
Actual \ominus	10	40	50
	40	60	100

		Predicted class	
		P	N
Actual Class	P	True Positives (TP)	False Negatives (FN)
	N	False Positives (FP)	True Negatives (TN)

$$ERR = \frac{FP + FN}{FP + FN + TP + TN} = 1 - ACC$$

$$ACC = \frac{TP + TN}{FP + FN + TP + TN} = 1 - ERR$$

$$FPR = \frac{FP}{N} = \frac{FP}{FP + TN}$$

$$TPR = \frac{TP}{P} = \frac{TP}{FN + TP}$$

$$PRE = \frac{TP}{TP + FP}$$

$$REC = TPR = \frac{TP}{P} = \frac{TP}{FN + TP}$$

$$F_1 = 2 \cdot \frac{PRE \cdot REC}{PRE + REC}$$

```
[1] 1 import numpy as np
    2 import matplotlib.pyplot as plt
    3 from sklearn.metrics import confusion_matrix
    4 from sklearn.metrics import classification_report
    5 from sklearn.metrics import ConfusionMatrixDisplay
    6 #from sklearn import metrics
```

	Predicted \oplus		Predicted \ominus	
Actual \oplus	30	20	50	
	10	40	50	
	40	60	100	

```
[2] 1 y_pred = np.array([1,1,1,1,0,0,0,0,0,0])
    2 y_true = np.array([1,1,1,0,1,1,0,0,0,0])
```


```
[3] 1 print(confusion_matrix(y_true,y_pred)) # the default sorts entries, so it shows 0 then 1
```

```
[[4 1]
 [2 3]]
```

```
[4] 1 print(confusion_matrix(y_true,y_pred,labels=[1,0])) # to reverse this like ppt slide add labels argument
```

```
[[3 2]
 [1 4]]
```

```
1 ConfusionMatrixDisplay.from_predictions(y_true, y_pred, labels=[1,0],cmap='Blues')
2 plt.show()
```



	<i>Predicted</i> ⊕	<i>Predicted</i> ⊖	
<i>Actual</i> ⊕	30	20	50
<i>Actual</i> ⊖	10	40	50
	40	60	100

```
[7] 1 print(classification_report(y_true, y_pred))
```

	precision	recall	f1-score	support
0	0.67	0.80	0.73	5
1	0.75	0.60	0.67	5
accuracy			0.70	10
macro avg	0.71	0.70	0.70	10
weighted avg	0.71	0.70	0.70	10



```
1 from sklearn.metrics import precision_score  
2 precision_score(y_true, y_pred)
```

0.75

	Predicted \oplus	Predicted \ominus	
Actual \oplus	30	20	50
Actual \ominus	10	40	50
	40	60	100

- In binary case, metrics like precision, recall, and f1-score are defined for positive case only
- Averages can be used for multiclass cases

sklearn.metrics.confusion_matrix

`sklearn.metrics.confusion_matrix(y_true, y_pred, *, labels=None, sample_weight=None, normalize=None)`

[\[source\]](#)

Compute confusion matrix to evaluate the accuracy of a classification.

By definition a confusion matrix C is such that $C_{i,j}$ is equal to the number of observations known to be in group i and predicted to be in group j .

Thus in binary classification, the count of true negatives is $C_{0,0}$, false negatives is $C_{1,0}$, true positives is $C_{1,1}$ and false positives is $C_{0,1}$.

Read more in the [User Guide](#).

Parameters: **y_true** : *array-like of shape (n_samples,)*

Ground truth (correct) target values.

y_pred : *array-like of shape (n_samples,)*

Estimated targets as returned by a classifier.

labels : *array-like of shape (n_classes), default=None*

List of labels to index the matrix. This may be used to reorder or select a subset of labels. If `None` is given, those that appear at least once in `y_true` or `y_pred` are used in sorted order.

sample_weight : *array-like of shape (n_samples,), default=None*

Sample weights.

New in version 0.18.

normalize : *{'true', 'pred', 'all'}, default=None*

Normalizes confusion matrix over the true (rows), predicted (columns) conditions or all the population. If `None`, confusion matrix will not be normalized.

```
>>> from sklearn.metrics import confusion_matrix
>>> y_true = [2, 0, 2, 2, 0, 1]
>>> y_pred = [0, 0, 2, 2, 0, 2]
>>> confusion_matrix(y_true, y_pred)
array([[2, 0, 0],
       [0, 0, 1],
       [1, 0, 2]])
```

```
>>> y_true = ["cat", "ant", "cat", "cat", "ant", "bird"]
>>> y_pred = ["ant", "ant", "cat", "cat", "ant", "cat"]
>>> confusion_matrix(y_true, y_pred, labels=["ant", "bird", "cat"])
array([[2, 0, 0],
       [0, 0, 1],
       [1, 0, 2]])
```

```
>>> from sklearn.metrics import classification_report
>>> y_true = [0, 1, 2, 2, 2]
>>> y_pred = [0, 0, 2, 2, 1]
>>> target_names = ['class 0', 'class 1', 'class 2']
>>> print(classification_report(y_true, y_pred, target_names=target_names))
```

	precision	recall	f1-score	support
class 0	0.50	1.00	0.67	1
class 1	0.00	0.00	0.00	1
class 2	1.00	0.67	0.80	3
accuracy			0.60	5
macro avg	0.50	0.56	0.49	5
weighted avg	0.70	0.60	0.61	5

Overall accuracy = 0.60

Note: in the binary case, the precision and recall for the problem are the values for the positive class

Summary

Know the following:

- Accuracy / error rate
- TP, FP, TN, FN in confusion matrix
- Precision
- Recall
- F1 Score

Don't need to memorize, just be familiar with the concepts and know how to look up definition as needed...

Evaluating Classifiers

MORE PROBLEMS

Chapter 2.1

Difficulty: 2

Given that:

Total = 100

False Negatives = 10

Precision = $\frac{4}{5}$

Recall = $\frac{6}{7}$

Can you complete the contingency matrix.

	Predicted +	Predicted -	
Actual +			
Actual -			

Chapter 2.1

Difficulty: 4

Two binary classifiers are used to predicted whether a patent has a life threatening diseases or not. Decide whether Classifier A or B would be better at reducing casualties.

A

10	10	20
40	9940	9980
50	9950	10000

B

13	7	20
87	9813	9980
100	9900	10000

Two contingency matrices.